

[Das...](#) / [My...](#) / [Computer Eng...](#) / [CEIT-even-...](#) / [OS-even-...](#) / [Theory: rand...](#) / [Random Quiz 4 : Scheduling, signals, segmentation,...](#)

<b>Started on</b>	Thursday, 16 February 2023, 9:01 PM
<b>State</b>	Finished
<b>Completed on</b>	Thursday, 16 February 2023, 10:08 PM
<b>Time taken</b>	1 hour 6 mins
<b>Grade</b>	13.65 out of 15.00 (91.01%)

## Question 1

Correct

Mark 1.00 out of 1.00

Which of the following statements is false ?

Select one:

- ☐ a. Response time is more predictable in preemptive systems than in non preemptive systems.
- ☒ b. Real time systems generally use non preemptive CPU scheduling. ✓
- ☐ c. Time sharing systems generally use preemptive CPU scheduling.
- ☐ d. A process scheduling algorithm is preemptive if the CPU can be forcibly removed from a process.

Your answer is correct.

The correct answer is: Real time systems generally use non preemptive CPU scheduling.

## Question 2

Correct

Mark 1.00 out of 1.00

Mark whether the concept is related to scheduling or not.

Yes	No		
<input checked="" type="radio"/>	<input type="radio"/>	ready-queue	✓
<input type="radio"/>	<input checked="" type="radio"/>	file-table	✓
<input checked="" type="radio"/>	<input type="radio"/>	timer interrupt	✓
<input checked="" type="radio"/>	<input type="radio"/>	context-switch	✓
<input checked="" type="radio"/>	<input type="radio"/>	runnable process	✓

ready-queue: Yes

file-table: No

timer interrupt: Yes

context-switch: Yes

runnable process: Yes

## Question 3

Correct

Mark 1.00 out of 1.00

Map each signal with it's meaning

SIGALRM	Timer Signal from alarm()	✓
SIGPIPE	Broken Pipe	✓
SIGCHLD	Child Stopped or Terminated	✓
SIGSEGV	Invalid Memory Reference	✓
SIGUSR1	User Defined Signal	✓

The correct answer is: SIGALRM → Timer Signal from alarm(), SIGPIPE → Broken Pipe, SIGCHLD → Child Stopped or Terminated, SIGSEGV → Invalid Memory Reference, SIGUSR1 → User Defined Signal

## Question 4

Correct

Mark 1.00 out of 1.00

Order the sequence of events, in scheduling process P1 after process P0

context of P0 is saved in P0's PCB	3	✓
context of P1 is loaded from P1's PCB	4	✓
Control is passed to P1	5	✓
timer interrupt occurs	2	✓
Process P1 is running	6	✓
Process P0 is running	1	✓

Your answer is correct.

The correct answer is: context of P0 is saved in P0's PCB → 3, context of P1 is loaded from P1's PCB → 4, Control is passed to P1 → 5, timer interrupt occurs → 2, Process P1 is running → 6, Process P0 is running → 1

## Question 5

Correct

Mark 1.00 out of 1.00

Match the names of PCB structures with kernel

xv6  ✓linux  ✓

The correct answer is: xv6 → struct proc, linux → struct task\_struct

## Question 6

Partially correct

Mark 0.75 out of 1.00

Select all the correct statements about signals

Select one or more:

- ☐ a. Signal handlers once replaced can't be restored
- ☐ b. The signal handler code runs in kernel mode of CPU
- ☐ c. SIGKILL definitely kills a process because it's code runs in kernel mode of CPU
- ☒ d. Signals are delivered to a process by another process ✗
- ☒ e. SIGKILL definitely kills a process because it can't be caught or ignored, and it's default action terminates the process ✓
- ☒ f. Signals are delivered to a process by kernel ✓
- ☒ g. The signal handler code runs in user mode of CPU ✓
- ☒ h. A signal handler can be invoked asynchronously or synchronously depending on signal type ✓

Your answer is partially correct.

You have selected too many options.

The correct answers are: Signals are delivered to a process by kernel, A signal handler can be invoked asynchronously or synchronously depending on signal type, The signal handler code runs in user mode of CPU, SIGKILL definitely kills a process because it can't be caught or ignored, and it's default action terminates the process

## Question 7

Correct

Mark 1.00 out of 1.00

Select the state that is not possible after the given state, for a process:

New:  ✓Ready:  ✓Running:  ✓Waiting:  ✓

## Question 8

Correct

Mark 1.00 out of 1.00

Which of the following are NOT a part of job of a typical compiler?

- ☐ a. Process the # directives in a C program
- ☒ b. Check the program for logical errors ✓
- ☐ c. Check the program for syntactical errors
- ☒ d. Suggest alternative pieces of code that can be written ✓
- ☐ e. Invoke the linker to link the function calls with their code, extern globals with their declaration
- ☐ f. Convert high level language code to machine code

Your answer is correct.

The correct answers are: Check the program for logical errors, Suggest alternative pieces of code that can be written

## Question 9

Correct

Mark 1.00 out of 1.00

Select the compiler's view of the process's address space, for each of the following MMU schemes:  
(Assume that each scheme, e.g. paging/segmentation/etc is effectively utilised)

Paging	one continuous chunk	✓
Segmentation	many continuous chunks of variable size	✓
Relocation + Limit	one continuous chunk	✓
Segmentation, then paging	many continuous chunks of variable size	✓

Your answer is correct.

The correct answer is: Paging → one continuous chunk, Segmentation → many continuous chunks of variable size, Relocation + Limit → one continuous chunk, Segmentation, then paging → many continuous chunks of variable size

Question **10**

Partially correct

Mark 0.67 out of 1.00

Which of the following parts of a C program do not have any corresponding machine code ?

- ☒ a. `#directives` ✓
- ☐ b. function calls
- ☐ c. global variables
- ☒ d. `typedefs` ✓
- ☐ e. expressions
- ☐ f. pointer dereference
- ☐ g. local variable declaration

Your answer is partially correct.

You have correctly selected 2.

The correct answers are: `#directives`, `typedefs`, global variables

## Question 11

Partially correct

Mark 1.38 out of 2.00

Select all the correct statements about the state of a process.

- ☒ a. A process changes from running to ready state on a timer interrupt ✓
- ☐ b. A process in ready state is ready to receive interrupts
- ☒ c. A running process may terminate, or go to wait or become ready again ✓
- ☒ d. Processes in the ready queue are in the ready state ✓
- ☒ e. Typically, it's represented as a number in the PCB ✓
- ☒ f. A process in ready state is ready to be scheduled ✓
- ☐ g. A waiting process starts running after the wait is over
- ☒ h. It is not maintained in the data structures by kernel, it is only for conceptual understanding of programmers ✗
- ☒ i. Changing from running state to waiting state results in "giving up the CPU" ✓
- ☐ j. A process waiting for any condition is woken up by another process only
- ☐ k. A process changes from running to ready state on a timer interrupt or any I/O wait
- ☐ l. A process can self-terminate only when it's running
- ☒ m. A process that is running is not on the ready queue ✓
- ☒ n. A process waiting for I/O completion is typically woken up by the particular interrupt handler code ✓

Your answer is partially correct.

You have correctly selected 8.

The correct answers are: Typically, it's represented as a number in the PCB, A process in ready state is ready to be scheduled, Processes in the ready queue are in the ready state, A process that is running is not on the ready queue, A running process may terminate, or go to wait or become ready again, A process changes from running to ready state on a timer interrupt, Changing from running state to waiting state results in "giving up the CPU", A process can self-terminate only when it's running, A process waiting for I/O completion is typically woken up by the particular interrupt handler code

Question **12**

Correct

Mark 1.00 out of 1.00

Select all the correct statements about zombie processes

Select one or more:

- ☒ a. A process becomes zombie when it finishes, and remains zombie until parent calls wait() on it ✓
- ☐ b. A process becomes zombie when it's parent finishes
- ☐ c. Zombie processes are harmless even if OS is up for long time
- ☒ d. If the parent of a process finishes, before the process itself, then after finishing the process is typically attached to 'init' as parent ✓
- ☒ e. A zombie process occupies space in OS data structures ✓
- ☒ f. A process can become zombie if it finishes, but the parent has finished before it ✓
- ☐ g. A zombie process remains zombie forever, as there is no way to clean it up
- ☒ h. init() typically keeps calling wait() for zombie processes to get cleaned up ✓

Your answer is correct.

The correct answers are: A process becomes zombie when it finishes, and remains zombie until parent calls wait() on it, A process can become zombie if it finishes, but the parent has finished before it, A zombie process occupies space in OS data structures, If the parent of a process finishes, before the process itself, then after finishing the process is typically attached to 'init' as parent, init() typically keeps calling wait() for zombie processes to get cleaned up


## Question 13

Correct

Mark 1.00 out of 1.00

Mark statements True/False w.r.t. change of states of a process. Note that a statement is true only if the claim and argument both are true.

Reference: The process state diagram (and your understanding of how kernel code works). Note - the diagram does not show zombie state!



**Figure 3.2** Diagram of process state.

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	A process in WAITING state can not become RUNNING because the event it's waiting for has not occurred and it has not been moved to ready queue yet	✓
<input type="radio"/>	<input checked="" type="radio"/>	A process in READY state can not go to WAITING state because the resource on which it will WAIT will not be in use when process is in READY state.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Only a process in READY state is considered by scheduler	✓
<input checked="" type="radio"/>	<input type="radio"/>	Every forked process has to go through ZOMBIE state, at least for a small duration.	✓
<input type="radio"/>	<input checked="" type="radio"/>	A process only in RUNNING state can become TERMINATED because scheduler moves it to ZOMBIE state first	✓

A process in WAITING state can not become RUNNING because the event it's waiting for has not occurred and it has not been moved to ready queue yet: True

A process in READY state can not go to WAITING state because the resource on which it will WAIT will not be in use when process is in READY state.: False

Only a process in READY state is considered by scheduler: True

Every forked process has to go through ZOMBIE state, at least for a small duration.: True

A process only in RUNNING state can become TERMINATED because scheduler moves it to ZOMBIE state first: False



Question **14**

Partially correct

Mark 0.86 out of 1.00

Mark True/False

Statements about scheduling and scheduling algorithms

True	False		
<input type="radio"/>	<input checked="" type="radio"/>	Generally the voluntary context switches are much more than non-voluntary context switches on a Linux system.	✗
<input checked="" type="radio"/>	<input type="radio"/>	A scheduling algorithm is non-preemptive if it does context switch only if a process voluntarily relinquishes CPU or it terminates.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Processor Affinity refers to memory accesses of a process being stored on cache of that processor	✓
<input checked="" type="radio"/>	<input type="radio"/>	xv6 code does not care about Processor Affinity	✓
<input checked="" type="radio"/>	<input type="radio"/>	Statistical observations tell us that most processes have large number of small CPU bursts and relatively smaller numbers of large CPU bursts.	✓
<input type="radio"/>	<input checked="" type="radio"/>	On Linuxes the CPU utilisation is measured as the time spent in scheduling the idle thread	✓ It's the negation of this. Time NOT spent in idle thread.
<input checked="" type="radio"/>	<input type="radio"/>	Response time will be quite poor on non-interruptible kernels	✓

Generally the voluntary context switches are much more than non-voluntary context switches on a Linux system.: True

A scheduling algorithm is non-preemptive if it does context switch only if a process voluntarily relinquishes CPU or it terminates.: True

Processor Affinity refers to memory accesses of a process being stored on cache of that processor: True

xv6 code does not care about Processor Affinity: True

Statistical observations tell us that most processes have large number of small CPU bursts and relatively smaller numbers of large CPU bursts.: True

On Linuxes the CPU utilisation is measured as the time spent in scheduling the idle thread: False

Response time will be quite poor on non-interruptible kernels: True

[◀ Random Quiz - 3 \(processes, memory management, event driven kernel\), compilation-linking-loading, ipc-pipes](#)

Jump to...

[Homework questions: Basics of MM, xv6 booting ▶](#)